Attention Is All You Need
A 50-page educational overview inspired by the 2017 Transformer paper.

This document summarizes concepts related to the Transformer architecture, self-attention, training dynamics, applications, and historical impact.

Page 1: Origins of sequence modeling

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Origins of sequence modeling. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 2: Problems with RNNs

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Problems with RNNs. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 3: Encoder-decoder architecture

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Encoder-decoder architecture. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 4: Self-attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Self-attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 5: Scaled dot-product attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Scaled dot-product attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 6: Multi-head attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Multi-head attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 7: Positional encoding

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Positional encoding. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 8: Residual connections

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Residual connections. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 9: Layer normalization

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Layer normalization. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 10: Feed-forward networks

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Feed-forward networks. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 11: Training transformers

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Training transformers. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 12: Complexity analysis

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Complexity analysis. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 13: Machine translation

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Machine translation. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 14: Transformer variants

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Transformer variants. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 15: BERT

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: BERT. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 16: GPT

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: GPT. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 17: Vision transformers

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Vision transformers. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 18: Long-context models

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Long-context models. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 19: Attention visualization

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Attention visualization. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 20: Applications in NLP

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Applications in NLP. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 21: Scaling laws

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Scaling laws. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 22: Limitations

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Limitations. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 23: Hallucinations

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Hallucinations. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 24: Efficient attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Efficient attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 25: Sparse attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Sparse attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 26: Mixture of experts

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Mixture of experts. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 27: Multimodal models

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Multimodal models. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 28: Future directions

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Future directions. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 29: Origins of sequence modeling

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Origins of sequence modeling. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 30: Problems with RNNs

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Problems with RNNs. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 31: Encoder-decoder architecture

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Encoder-decoder architecture. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 32: Self-attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Self-attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 33: Scaled dot-product attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Scaled dot-product attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 34: Multi-head attention

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Multi-head attention. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 35: Positional encoding

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Positional encoding. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 36: Residual connections

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Residual connections. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 37: Layer normalization

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Layer normalization. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 38: Feed-forward networks

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Feed-forward networks. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 39: Training transformers

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Training transformers. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 40: Complexity analysis

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Complexity analysis. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 41: Machine translation

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Machine translation. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 42: Transformer variants

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Transformer variants. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 43: BERT

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: BERT. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 44: GPT

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: GPT. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 45: Vision transformers

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Vision transformers. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 46: Long-context models

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Long-context models. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 47: Attention visualization

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Attention visualization. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 48: Applications in NLP

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Applications in NLP. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 49: Scaling laws

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Scaling laws. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.

Page 50: Limitations

The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. The paper 'Attention Is All You Need' introduced the Transformer architecture, replacing recurrence with attention mechanisms. This enabled greater parallelization, improved long-range dependency handling, and reshaped modern AI systems. Self-attention computes relationships among tokens in a sequence, producing contextual representations. The architecture later influenced large language models, vision systems, and multimodal AI. This page focuses on: Limitations. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries. Transformer research expanded rapidly after 2017, leading to foundation models used across industries.

Key takeaway: attention mechanisms allow models to weigh different parts of input sequences dynamically.